

Introduction to automation using Nextflow:

A tutorial through examples

Phelelani Mpangase

Sydney Brenner Institute for Molecular Bioscience
University of the Witwatersrand
Johannesburg
South Africa

Outline

- 1 Introduction to Nextflow
 - Introduction
 - Nextflow Script
 - Partial Execution
 - Visualising the Workflow
- 2 Generalising and Extending
 - Parameters
 - Channels
 - Generalising Our Example
 - Managing Grouped Files
 - On absolute paths
- 3 Nextflow and Docker
 - Docker & Singularity Containers
 - Directory & File Access
- 4 Executors
 - Executors
 - Nextflow on a cluster (HPC)
 - Scheduler + Docker
 - Amazon EC2
- 5 How it All Fits

Introduction to Nextflow

Introduction to Nextflow

Introduction

Resources

- <https://github.com/phelelani/nf-tut-2023>

Workflow Languages

Many scientific applications require

- Multiple data files
- Multiple applications
- Perhaps different parameters

General purpose languages not well suited

- Too low a level of abstraction
- Does not separate workflow from application
- Not reproducible

Workflow Languages

Many scientific applications require

- Multiple data files
- Multiple applications
- Perhaps different parameters

General purpose languages not well suited

- Too low a level of abstraction
- Does not separate workflow from application
- Not reproducible

Nextflow enables reproducible computational workflows

To the Editor:

The increasing complexity of readouts for omics analyses goes hand-in-hand with concerns about the reproducibility of experiments that analyze 'big data'¹⁻³. When analyzing very large data sets, the main source of computational irreproducibility arises from a lack of good practice pertaining to software and database usage⁴⁻⁶. Small variations across computational platforms also contribute to computational irreproducibility by producing numerical instability⁷, which is especially relevant to high-performance computational (HPC) environments that are routinely used for omics analyses⁸. We present a solution to this instability named Nextflow, a workflow management system that uses Docker technology for the multi-scale handling of containerized computation.

In silico workflow management systems are an integral part of large-scale biological analyses. These systems enable the rapid prototyping and deployment of pipelines that combine complementary software packages. In genomics the simplest pipelines, such as Kallisto and Sleuth⁹, combine an RNA-seq quantification method with a differential expression module (**Supplementary Fig. 1**). Complexity rapidly increases when all aspects of a given analysis are included. For example,

the Sanger Companion pipeline¹⁰ bundles 39 independent software tools and libraries into a genome annotation suite. Handling such a large number of software packages, some of which may be incompatible, is a challenge. The conflicting requirements of frequent software updates and maintaining the reproducibility of original results provide another unwelcome wrinkle. Together with these problems, high-throughput usage of complex pipelines can also be burdened by the hundreds of intermediate files often produced by individual tools. Hardware fluctuations in these types of pipelines, combined with poor error handling, could result in considerable readout instability.

Nextflow (<http://nextflow.io>; **Supplementary Methods, Supplementary Note and Supplementary Code 1**) is designed to address numerical instability, efficient parallel execution, error tolerance, execution provenance and traceability. It is a domain-specific language that enables rapid pipeline development through the adaptation of existing pipelines written in any scripting language.

We present a qualitative comparison between Nextflow and other similar tools in **Table 1** (ref. 11). We found that multi-scale containerization, which makes it possible to

Nextflow

Groovy-based language

- Expressing workflows
- Portable
 - works on most Unix-like systems
- Very easy to install
 - NB: requires Java 7, 8
- Scalable
- Supports Docker/Singularity
- Supports a range of scheduling systems

Nextflow

Groovy-based language

- Expressing workflows
- Portable
 - works on most Unix-like systems
- Very easy to install
 - NB: requires Java 7, 8
- Scalable
- Supports Docker/Singularity
- Supports a range of scheduling systems

Key concepts of Nextflow

- **Processes:**
 - actual work being done (usually simple).
 - call program that does the analysis.
- **Channels:**
 - for communication between processes.
 - handles inputs and outputs.
- When all inputs ready, process is executed.
- Each process runs in its own directory (files are staged).
- Supports resumption of previous partial runs.

Introduction to Nextflow

Nextflow Script

Exercise 1

You have an input file with 6 columns (see below), where column 2 is an "index" column. Identify rows that have identical indexes (column 2) and remove them from the file.

Your input file looks like this:

11	11:189256	0	189256	A	G
11	11:193788	0	193788	T	C
11	11:194062	0	194062	T	C
11	11:194228	0	194228	A	G
11	11:193788	0	193788	A	C

Simple Example: Using BASH

Input is a file

- With 6 columns
- Column 2 is an index column
- Identify rows with identical field 2
- Remove identical rows

```
11  11:189256  0  189256  A  G
11  11:193788  0  193788  T  C
11  11:194062  0  194062  T  C
11  11:194228  0  194228  A  G
11  11:193788  0  193788  A  C
```

Using BASH:

```
cut -f 2 data/11.bim | sort | uniq -d > dups
grep -v -f dups data/11.bim > 11.clean
```

Simple Example: Using nextflow

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  input_ch = Channel.fromPath("data/11.bim")
5
6  process getIDs {
7      input:
8      path(input_ch)
9
10     output:
11     path("ids"), emit: id_ch
12     path("11.bim"), emit: orig_ch
13
14     """
15     cut -f 2 ${input_ch} | sort > ids
16     """
17 }
18
19 process getDups {
20     input:
21     path(id_ch)
22
23     output:
24     path("dups"), emit: dups_ch
25
26     """
27     uniq -d ${id_ch} > dups
28     touch ignore
29     """
30 }
```

Simple Example: Using nextflow

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  input_ch = Channel.fromPath("data/11.bim")
5
6  process getIDs {
7      input:
8          path(input_ch)
9
10     output:
11         path("ids"), emit: id_ch
12         path("11.bim"), emit: orig_ch
13
14     """
15     cut -f 2 ${input_ch} | sort > ids
16     """
17 }
18
19 process getDups {
20     input:
21         path(id_ch)
22
23     output:
24         path("dups"), emit: dups_ch
25
26     """
27     uniq -d ${id_ch} > dups
28     touch ignore
29     """
30 }
```

```
31 process removeDups {
32     input:
33         path(dups_ch)
34         path(orig_ch)
35
36     output:
37         path("clean.bim"), emit: output
38
39     """
40     grep -v -f ${dups_ch} ${orig_ch} > clean.bim
41     """
42 }
43
44 workflow {
45     getIDs(input_ch)
46     getDups(getIDs.out.id_ch)
47     removeDups(getDups.out.dups_ch, getIDs.out.orig_ch).
48         subscribe { print "Done!" }
49 }
```

Simple Example: Using nextflow

```
$ nextflow run cleandups.nf
```

```
N E X T F L O W ~ version 19.04.1
Launching `cleandups.nf` [soggy_jennings] - revision: 795e2aa39d
[warm up] executor > local
executor > local (3)
[84/7e1ad1] process > getIDs      [100%] 1 of 1 ▢
[19/cc8bf9] process > getDups     [100%] 1 of 1 ▢
[f9/ed086d] process > removeDups [100%] 1 of 1 ▢
Completed at: 31-Jul-2019 09:00:50
Duration      : 1.5s
CPU hours     : (a few seconds)
Succeeded    : 3
```

```
--work
| |--90
| | |--cebf3649d883f88381e32b4912b560
| | | |--ids -> /Users/phele/day4/work/b3/aa0380f2a1bca447259b7ffd390083/ids
| | | |--ignore
| |--9c
| | |--e0cb7d8d26682d7d4a1c44392f2bb3
| | | |--11.bim -> /Users/phele/day4/data/11.bim
| | | |--clean.bim
| | | |--dups -> /Users/phele/day4/work/90/cebf3649d883f88381e32b4912b560/dups
| |--b3
| | |--aa0380f2a1bca447259b7ffd390083
| | | |--11.bim -> /Users/phele/day4/data/11.bim
| | | |--ids
```

Exercise 2

Change the script so that you use `stdin` or `stdout` in the `getIDs` and `getDups` processes to avoid the use of the temporary file ids. You can see the solution: [ex2-cleandups-stdin.nf](https://github.com/nextflow-io/nextflow/blob/master/examples/ex2-cleandups-stdin.nf)!.

Introduction to Nextflow

Partial Execution

Partial Execution

If execution of workflow is only partial

- Because of error
- Only need to resume from process that failed

```
nextflow run cleandups.nf -resume
```

Introduction to Nextflow


Visualising the Workflow

Visualising the Workflow

Nextflow supports several visualisation tools:

-with-dag

```
nextflow run cleandups.nf -with-dag <file-name>
```



Visualising the Workflow

Nextflow supports several visualisation tools:

-with-dag

```
nextflow run cleandups.nf -with-dag <file-name>
```

-with-timeline

```
nextflow run cleandups.nf -with-timeline <file-name>
```

Processes execution timeline

Launch time: 05 Jun 2018 10:41
Elapsed time: 2.9s



Created with Nextflow -- <http://nextflow.io>

Visualising the Workflow

Nextflow supports several visualisation tools:

-with-dag

```
nextflow run cleandups.nf -with-dag <file-name>
```

-with-timeline

```
nextflow run cleandups.nf -with-timeline <file-name>
```

-with-report

```
nextflow run cleandups.nf -with-report <filename>
```

The screenshot displays a 'Nextflow Report' interface. At the top, there are tabs for 'Summary', 'Resources', and 'Tasks', and a user identifier '[evil_keller]'. The main heading is 'Nextflow workflow report' followed by '[evil_keller]'. A green status bar indicates 'Workflow execution completed successfully!'. Below this, the 'Run times' section shows the execution period from 'Tue Jun 05 10:41:20 SAST 2018' to 'Tue Jun 05 10:41:22 SAST 2018', noting it was completed 5 minutes ago with a duration of 1.8s. A green progress bar shows '3 succeeded'. The 'Nextflow command' section contains the command: 'nextflow run cleandups.nf -with-timeline timeline.html -with-dag timeline.dot -with-report report.html'. At the bottom, a table lists various metadata items:

CPU-Hours	(a few seconds)
Launch directory	/home/phelani/2018_courses/nextflow-course
Work directory	/home/phelani/2018_courses/nextflow-course/work
Project directory	/home/phelani/2018_courses/nextflow-course
Script name	cleandups.nf
Script ID	705c0ae79d5c8fa2d2961d9f9448c1e
Workflow session	68916148-f943-4369-bca5-7cc8e5d296e3
Workflow profile	standard
Nextflow version	version 0.29.1, build 4804 (10-05-2018 07:47 UTC)

Generalising and Extending

Extending the Example

- Parameterise the input
- Want output to go to convenient place
- Workflow takes in multiple input files – processes are executed on each in turn.
- Complication : may need to carry the base name of the input to the final output;
- Can repeat some steps for different parameters.

Generalising and Extending

Parameters

Parameters

In Nextflow file:

```
input_ch = Channel.fromPath(params.data)
```

And run it like this

```
nextflow run phylo1.nf --data data/polyseqs.fa
```

Generalising and Extending

Channels

Data Types in Channels

Channels support different types:

- path
- stdin
- env
- tuple

Many, many operations you can do on channels and their contents

bind	buffer	close
filter	map/reduce	group
join, merge	mix	copy
split	spread	fork
count	min/max/sum	print/view

Creating Channels

```
Channel.create()
Channel.empty
Channel.from("blast", "plink")
Channel.fromPath("data/*.fa")
Channel.fromFilePairs("data/{YRI,CEU,BEB}.*")
Channel.watchPath("*.fa")
```

Generalising and Extending

Generalising Our Example

Workflow: Multiple Inputs

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.data_dir = "data"
5  input_ch = Channel.fromPath("${params.data_dir}/*.bim")
6
7  process getIDs {
8      input:
9      path(input)
10
11      output:
12      path("${input.baseName}.ids"), emit: id_ch
13      path("${input}"), emit: orig_ch
14
15      """
16      cut -f 2 ${input} | sort > ${input.baseName}.ids
17      """
18  }
19
20  process getDups {
21      input:
22      path(input)
23
24      output:
25      path("${input.baseName}.dups"), emit: dups_ch
26
27      """
28      uniq -d ${input} > "${input.baseName}.dups"
29      touch ignore
30      """
31  }
```

Workflow: Multiple Inputs

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.data_dir = "data"
5  input_ch = Channel.fromPath("${params.data_dir}/*.bim")
6
7  process getIDs {
8      input:
9          path(input)
10
11      output:
12          path("${input.baseName}.ids"), emit: id_ch
13          path("${input}"), emit: orig_ch
14
15      """
16      cut -f 2 ${input} | sort > ${input.baseName}.ids
17      """
18  }
19
20  process getDups {
21      input:
22          path(input)
23
24      output:
25          path("${input.baseName}.dups"), emit: dups_ch
26
27      """
28      uniq -d ${input} > "${input.baseName}.dups"
29      touch ignore
30      """
31  }
```

```
32  process removeDups {
33      publishDir "output", pattern: "${badids.baseName}.bim",
34          overwrite:true, mode:'copy'
35
36      input:
37          path(badids)
38          path(orig)
39
40      output:
41          path("${badids.baseName}_clean.bim"), emit: cleaned_ch
42
43      """
44      grep -v -f ${badids} ${orig} > ${badids.baseName}_clean.bim
45      """
46  }
47
48  workflow {
49      getIDs(input_ch)
50      getDups(getIDs.out.id_ch)
51      removeDups(getDups.out.dups_ch, getIDs.out.orig_ch)
52  }
```

Workflow: Multiple Inputs

```
$ nextflow run cleandups.nf

Launching `cleandups.nf` [distracted_hodgkin] - revision: 29fdb384a6
[warm up] executor > local
executor > local (9)
[1a/431eb7] process > getIDs      [100%] 3 of 3 ▮
[cc/fc0aaa] process > getDups     [100%] 3 of 3 ▮
[03/c31154] process > removeDups [100%] 3 of 3 ▮
Completed at: 31-Jul-2019 10:26:23
Duration      : 2s
CPU hours     : (a few seconds)
Succeeded     : 9
```


Exercise 4

Now try adding a process to our Nextflow example and for splitting the file but using different split values (solution: [ex4-cleandups-multi-params.nf](#)).

Workflow: Multiple Parameters

Now try splitting the file but use different split values

```
split -l 400 data.txt dataX
```

will produce files dataXaa, dataXab, dataXac and so on ...

Try:

```
1 splits = [400,500,600]
2
3 process splitIDs {
4   input:
5     path(bim)
6   each split
7
8   output:
9     path("*-${split-*}"), emit: output_ch
10
11   """
12   split -l ${split} ${bim} ${bim.baseName}-${split}-
13   """
14 }
```

Have a look at the modified Nextflow script: [ex4-cleandups-multi-params-mod.nf](#).

Generalising and Extending

Managing Grouped Files

Grouped Files

Use **PLINK** as an example.

```
## Short version of the command
plink --bfile /path/YRI --freq --out /tmp/YRI

## Long version of the command
plink --bed YRI.bed \
  --bim YRI.bim \
  --fam YRI.fam \
  --freq \
  --out /tmp/YRI
```

Problem:

- Pass the files on another channel(s) to be staged
- Pass the base name as value/or work it out

Pros/Cons

- Simple
- Need extra channel/some gymnastics

Grouped Files

Use **PLINK** as an example.

```
## Short version of the command
plink --bfile /path/YRI --freq --out /tmp/YRI

## Long version of the command
plink --bed YRI.bed \
  --bim YRI.bim \
  --fam YRI.fam \
  --freq \
  --out /tmp/YRI
```

Problem:

- Pass the files on another channel(s) to be staged
- Pass the base name as value/or work it out

Pros/Cons

- Simple
- Need extra channel/some gymnastics

RECAP CLOSURES

Simply, a *closure* is an anonymous function

- Code wrapped in braces {, }
- Default argument called *it*

```
[1,2,3].each { print it * it }
[1,2,3].each { num -> print num * num }
```

Grouped Files - Version 1: map

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.dir = "data/pops/"
5  dir = params.dir
6  params.pops = ["YRI", "CEU", "BEB"]
7
8  Channel
9      .from(params.pops)
10     .map { pop ->
11         [ file("${dir}/${pop}.bed"),
12           file("${dir}/${pop}.bim"),
13           file("${dir}/${pop}.fam") ]
14     }
15     .set { plink_data }
16
17 plink_data.subscribe { println "$it" }
```

Grouped Files - Version 1: map

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.dir = "data/pops/"
5  dir = params.dir
6  params.pops = ["YRI", "CEU", "BEB"]
7
8  Channel
9      .from(params.pops)
10     .map { pop ->
11         [ file("${dir}/${pop}.bed"),
12           file("${dir}/${pop}.bim"),
13           file("${dir}/${pop}.fam") ]
14     }
15     .set { plink_data }
16
17 plink_data.subscribe { println "$it" }
```

```
[data/pops/YRI.bed, data/pops/YRI.bim, data/pops/YRI.fam]
[data/pops/CEU.bed, data/pops/CEU.bim, data/pops/CEU.fam]
[data/pops/BEB.bed, data/pops/BEB.bim, data/pops/BEB.fam]
```

Grouped Files - Version 1: map

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.dir = "data/pops/"
5  dir = params.dir
6  params.pops = ["YRI", "CEU", "BEB"]
7
8  Channel
9      .from(params.pops)
10     .map { pop ->
11         [ file("${dir}/${pop}.bed"),
12           file("${dir}/${pop}.bim"),
13           file("${dir}/${pop}.fam") ]
14     }
15     .set { plink_data }
16
17 plink_data.subscribe { println "$it" }
```

```
16  process getFreq {
17      input:
18          tuple path(bed), path(bim), path(fam)
19
20      output:
21          path("${bed.baseName}.frq"), emit result
22
23      """
24      plink --bed $bed \
25            --bim $bim \
26            --fam $fam \
27            --freq \
28            --out "${bed.baseName}"
29      """
30  }
31
32  workflow {
33      getFreq(plink_data).view()
34  }
```

```
[data/pops/YRI.bed, data/pops/YRI.bim, data/pops/YRI.fam]
[data/pops/CEU.bed, data/pops/CEU.bim, data/pops/CEU.fam]
[data/pops/BEB.bed, data/pops/BEB.bim, data/pops/BEB.fam]
```


Grouped Files - Version 2: `fromFilePairs`

Use `fromFilePairs`.

- Takes a closure used to gather files together with the same key

```
x_ch = Channel.fromFilePairs( files ) { closure }
```

- Specify the files as a glob
- Closure associates each file with a key
- `fromPairs` puts all files with same key together
- Returns a list of pairs (key, list)

Grouped Files - Version 2: `fromFilePairs`

Use `fromFilePairs`.

- Takes a closure used to gather files together with the same key

```
x_ch = Channel.fromFilePairs( files ) { closure }
```

- Specify the files as a glob
- Closure associates each file with a key
- `fromPairs` puts all files with same key together
- Returns a list of pairs (key, list)

```
1  #!/usr/bin/env nextflow
2
3  commands = Channel.fromFilePairs("/usr/bin/*", size:-1) {
4      it.baseName[0]
5  }
6
7  commands.subscribe { k= it[0];
8      n=it[1].size();
9      println "There are $n files starting with $k";
10 }
```

A more complex example – default closure

```
1  Channel
2      .fromFilePairs
3      ("${params.dir}/*.{bed,fam,bim}",size:3, flat : true)
4      .ifEmpty { error "No matching plink files" }
5      .set { plink_data }
6
7  plink_data.subscribe { println "$it" }
```

Grouped Files - Version 2: fromFilePairs

Use `fromFilePairs`.

- Takes a closure used to gather files together with the same key

```
x_ch = Channel.fromFilePairs( files ) { closure }
```

- Specify the files as a glob
- Closure associates each file with a key
- `fromPairs` puts all files with same key together
- Returns a list of pairs (key, list)

```
[CEU, [data/pops/CEU.bed, data/pops/CEU.bim, data/pops/CEU.fam]]  
[YRI, [data/pops/YRI.bed, data/pops/YRI.bim, data/pops/YRI.fam]]  
[BEB, [data/pops/BEB.bed, data/pops/BEB.bim, data/pops/BEB.fam]]
```

```
1  #!/usr/bin/env nextflow  
2  
3  commands = Channel.fromFilePairs("/usr/bin/*", size:-1) {  
4      it.baseName[0]  
5  }  
6  
7  commands.subscribe { k= it[0];  
8      n=it[1].size();  
9      println "There are $n files starting with $k";  
10 }
```

A more complex example – default closure

```
1  Channel  
2      .fromFilePairs  
3      ("${params.dir}/*.{bed,fam,bim}",size:3, flat : true)  
4      .ifEmpty { error "No matching plink files" }  
5      .set { plink_data }  
6  
7  plink_data.subscribe { println "$it" }
```

Grouped Files - Version 2: fromFilePairs

```
1  process checkData {  
2      input:  
3      tuple val(pop), path(pl_files)  
4  
5      output:  
6      path("${pl_files[0]}.frq"), emit: result  
7  
8      """  
9      plink --bfile $base --freq --out pl_files[0].baseName  
10     """  
11 }
```

Grouped Files - Version 2: fromFilePairs

```
1 process checkData {  
2   input:  
3   tuple val(pop), path(pl_files)  
4  
5   output:  
6   path("${pl_files[0]}.frq"), emit: result  
7  
8   """  
9   plink --bfile $base --freq --out pl_files[0].baseName  
10  """  
11 }
```

```
1 process checkData {  
2   input:  
3   tuple val(pop), path(pl_files)  
4  
5   output:  
6   path("${pop}.frq"), emit: result  
7  
8   """  
9   plink --bfile $pop --freq --out $pop  
10  """  
11 }
```

Grouped Files - Final Version

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  params.dir = "data/pops/"
5  dir = params.dir
6  params.pops = ["YRI", "CEU", "BEB"]
7
8  Channel
9      .fromFilePairs("${params.dir}/{YRI,BEB,CEU}.{bed,bim,fam}",size:3) {
10      file -> file.baseName
11      }
12      .filter { key, files -> key in params.pops }
13      .set { plink_data }
14
15  process checkData {
16      input:
17      tuple val(pop), path(pl_files)
18
19      output:
20      path("${pop}.frq"), emit: result
21
22      """
23      plink --bfile $pop --freq --out $pop
24      """
25
26  workflow {
27      checkData(plink_data).view()
28  }
```

Exercise 5

Have a look at [ex5-weather.nf](#). In the data directory are set of data files for different years and months. First, I want you to use paste to combine all the files for the same year and month (paste joins files horizontal-wise). Then these new files should be concated.

Generalising and Extending

On absolute paths

Absolute paths

```
1 input = Channel.fromPath("/data/batch1/myfile.fa")
2
3 process show {
4     input:
5     path(data)
6
7     output:
8     path('see.out')
9
10    """
11    cp ${data} /home/scott/answer
12    """
13    ...
```

Nextflow and Docker

Nextflow and Docker

Docker & Singularity Containers

Docker & Singularity Containers

Light-weight virtualisation abstraction layer

- Currently runs on Unix like systems
 - Linux
 - macOS
- Windows support coming

Can create images locally or get from repositories

Docker

```
docker pull ubuntu
docker pull quay.io/banshee1221/h3agwas-plink
```

Singularity

```
singularity pull docker://ubuntu
singularity pull docker://quay.io/banshee1221/h3agwas-plink
```

- Docker/Singularity often run images in background
- Can also run interactively

Running Docker interactively

```
sudo docker run -t -i quay.io/banshee1221/h3agwas-plink
```

Running Singularity interactively

```
singularity shell docker://quay.io/banshee1221/h3agwas-plink
```

Running images

Docker

```
docker run <some-image-name>
```

Singularity

```
singularity exec <some-image-name>
```

Nextflow supports Docker & Singularity

- Well designed script should be highly portable
- Each process gets run as a separate image call
 - Under the hood, a **docker run** or a **singularity exec** is called
- Can use the same or different images for each process
 - Parameterisable

Assuming all processes use the same image:

```
## For Docker
nextflow run plink2.nf -with-docker quay.io/banshee1221/h3agwas-plink

## For Singularity
nextflow run plink.nf -with-singularity docker://quay.io/banshee1221/h3agwas-plink
```

Nextflow and Docker

Directory & File Access

Directory & File access

Nextflow Docker/Singularity support highly transparent – but pay attention to good practice

- For each process Docker/Singularity mounts the work directory for **that** process on the Docker/Singularity image.
- Files can be staged in and out using Nextflow mechanisms.
- Other files available: directories mounted through Docker/Singularity run time options or on the Docker image
- No other files on the host machine including the current directory
- Process executes in the Docker/Singularity environment

Directory & File access

```
1  #!/usr/bin/env nextflow
2  nextflow.enable.dsl=2
3
4  data = Channel.fromPath("data/pops/YRI.bim")
5
6  process see {
7      publishDir "count_out", overwrite:true, mode:'move'
8      echo true
9
10     input:
11     path(bim)
12
13     output:
14     path(count)
15
16     """
17     hostname
18     echo "Path is \$( pwd )\n "
19     echo "Parent directory has \$( ls .. )\n"
20     echo "My home directory has \$( ls /home/phele )\n"
21     wc -l ${bim} > count
22     ls
23     """
24 }
25
26 workflow {
27     see(data)
28 }
```

```
N E X T F L O W ~ version 0.21.2
Launching show_env.nf
[warm up] executor > local
[94/597f09] Submitted process > see (1)
89ad448ae0b2
Path is /home/scott/witsGWAS/dockerized/work/94/597f09ca6cc01c7be
Parent directory has 597f09ca6cc01c7be
My home directory has witsGWAS

YRI.bim
count
```


Directory & File access

Note that although the script's `pwd` shows:

```
/home/scott/witsGWAS/dockerized/work/94/597f09ca6cc01c7be
```

- Only these specific directories are mounted
- Only the files in the innermost directory are available

Any absolute paths (other than those used in staging) will result in error.

Profiles

In nextflow.config

```
1 profiles {  
2     ...  
3     docker {  
4         process.container = 'quay.io/banshee1221/h3agwas-plink:latest'  
5         docker.enabled = true  
6     }  
7 }
```

Now can run as:

```
nextflow run gwas.nf -profile docker
```

This can be extended in many ways

- Different processes can use different containers
- Can mount other host directories
- Can pass arbitrary Docker parameters

Executors

Executors

Executors

Executors

A Nextflow *executor* is the mechanism which Nextflow runs the code in each of the processes

- Default is `local`: process is run as a script

Many others

- PBS/Torque
- SLURM
- Amazon (AWS Batch)
- SGE (Sun Grid Engine)

Selecting an executor Annotating each process

- `executor` directive, e.g. `executor 'pbs'`
- resource constraints

Or, `nextflow.config` file

- either global or per-process

Executors

Nextflow on a cluster (HPC)

Running Nextflow on a cluster (HPC)

Script runs on the *head* node

- Nextflow uses the **executor** information to decide how the job should run
- Each process can be handled differently
- Nextflow submits each process to the job scheduler on your behalf (e.g, if using PBS/Torque, **qsub** is done)

Example

```
1 process {  
2   executor = 'pbs'  
3   queue = 'batch'  
4   scratch = true  
5   cpus = 5  
6   memory = '2GB'  
7 }
```

Executors

Scheduler + Docker

Scheduler + Docker

```
1 process.container = 'quay.io/banshee1221/h3agwas-plink:latest'
2 docker.enabled = false
3
4 process {
5     executor = 'pbs'
6     queue = 'batch'
7     scratch = true
8     cpus = 5
9     memory = '2GB'
10 }
```

Executors

Amazon EC2

Amazon EC2

Nextflow has native support for EC2

- You need an account on EC2
- Image (AMI) with the appropriate support

Launch your code:

```
nextflow cloud create GenomeCloud -c 5
```

If successful, Nextflow will give you the name of the headnode of your cluster

- **ssh** into it
- run Nextflow on it.

Afterwards shut down:

```
nextflow shutdown GenomeCloud
```

How it All Fits

Bioinformatics Workflows - Challenges & Solutions



Bioinformatics Workflows - Best Practices

